

COMP3048: Lecture 17

Code Optimisations

Matthew Pike

University of Nottingham, Ningbo

This Lecture: Optimization

- **Code improvement** or **optimization**:
what is it?
- High-level, intermediate-level, and low-level optimization.
- Time and space trade-offs.
- Specific optimizations; e.g.
 - Constant folding
 - Common subexpression evaluation
 - Inlining
- Interaction among Optimizations

Code Improvement (1)

The code generated by a compiler

- **must** be correct
(i.e., semantics-preserving translation)
- **should** also
 - run fast
 - be small
 - use as little space as possible

Code improvement is the process of improving the time and/or space behaviour of generated code **without** changing its functional behaviour; i.e. correctness **must** be preserved.

Code Improvement (2)

Consider:

```
w := 42;  
i := 0;  
while (i < 100) do begin  
    j := 0;  
    while (j < 200) do begin  
        x := (w * 10) * a[i];  
        y := y + x + b[j];  
        j := j + 1;  
    end;  
    i := i + 1;  
end
```

How might this code fragment be changed to (likely) make it run faster?

Code Improvement (3)

Example: Replacing the code fragment

$$f(x) + f(x)$$

by

$$2 * f(x)$$

saves a function call; likely reduces execution time.

Any caveat???

Only correct if f does **not** have any side effects!

Code Improvement (4)

Consider:

```
var x: Integer;  
...  
fun f (y: Integer): Integer =  
    begin  
        x = x + 1;  
        return x + y  
    end  
  
...  
x = 2;  
putint(f(2) + f(2))
```

This code fragment would print 11, whereas the result of printing $2 * f(2)$ would be 10.

Code Improvement (5)

Note: “Side effect” includes:

- Updates of variables, data-structures
- I/O and other changes to the system state
- Exceptions
- Non-termination

Optimization?

Code improvement usually referred to as “optimization”. However:

- Hardly ever possible to **guarantee** optimality under any mathematical measure.
- Not even always an improvement: not known what is going to happen at run-time, so “optimizing” for the **average expected** case.
- Careful and extensive **benchmarking** is often the only way to verify that an optimization indeed does improve generated code most of the time.

•
•
•

At What Level? (1)

Code improvement can be done at different levels:

- High level: source-to-source (AST) transformations.
- Intermediate level: transformations on intermediate representation, e.g.:
 - “bare-bones” high-level language
 - control/data flow graph representation
- Low level: transformations on machine code.

Each level suitable for different kinds of optimization.
Improve at all levels!

• • • • • • • • • • •

At What Level? (2)

Consider this code fragment:

```
if x then
    if y then
        printf(1)
    else
        printf(2)
else
    printf(3)
```

Anything that obviously could be improved at this level; i.e. the **source code** level?

At What Level? (3)

Resulting TAM code might be:

LOAD	[SB + 12]	#2:	LOADL	2
JUMPIFZ	#0		CALL	putint
LOADL	[SB + 13]	#3:	JUMP	#1
JUMPIFZ	#2	#0:	LOADL	3
LOADL	1		CALL	putint
CALL	putint	#1:		
JUMP	#3			

Now anything that could be improved;
i.e., at the *machine code* level?

•
•
•

At What Level? (4)

Information that was implicit in the high-level representation might become explicit at the intermediate level, thus enabling/facilitating certain optimizations.

Consider array indexing. High-level code fragments:

```
var x, y: array[1..100] Integer;  
...  
a := x[i] + y[i];
```

At What Level? (4)

Intermediate (C-like) code with explicit pointer arithmetic:

```
if (i < 1 || i > 100) then raise index_bounds;  
t1 := ^ (x + 4 * (i - 1));  
if (i < 1 || i > 100) then raise index_bounds;  
t2 := ^ (y + 4 * (i - 1));  
a := t1 + t2
```

(\wedge is the pointer dereferencing operator.)

•
•
•

At What Level? (5)

This could be optimized by reusing common subexpressions and eliminating redundant array bounds checks:

```
if (i < 1 || i > 100) then raise index_bounds;  
t0 := 4 * (i - 1)  
t1 := ^ (x + t0);  
t2 := ^ (y + t0);  
a := t1 + t2;
```

Time vs. Space (1)

Time and space optimizations are often in conflict.
Consider representing an array of Booleans:

- Each Boolean represented by one machine word:
 - fast access
 - wastes space.
- Each Boolean represented by a single bit:
 - space efficient
 - access requires extra operations (shifting and masking): takes time (and some *instruction* space)!

Time vs. Space (2)

In other cases, small is fast as well:

- Basic observation: accessing memory is slow. The fewer instructions and the fewer pieces of data, the fewer memory accesses, and the faster the execution.
- It is highly desirable to keep inner loops small so that they fit in the first-level **instruction** cache.
- It is desirable to keep the set of “currently accessed” memory locations small so that they fit in the first-level **data** cache.

•
•
•

Time vs. Space (3)

But then again, since memory access is very slow, avoiding a memory access could sometimes be worth a few extra instructions!

(Reason: Instruction fetching is typically much faster than data fetching because it is more predictable.)

Conclusion: the trade-off between time and space is a highly complicated issue!

In practice, one often have to make an educated guess, then verify by benchmarking.

Common Optimization Techniques

Applicable at the source-code (AST) level and/or intermediate level:

- Constant Folding
- Common Subexpression Elimination
- Algebraic Identities
- Copy Propagation
- Dead Code Elimination
- Strength Reduction
- Code Motion
- Loop Unrolling
- Inlining

Constant Folding (1)

Idea: evaluate (sub)expressions at compile-time where possible:

```
const pi: Double = 3.1416;  
var volume, radius: Double;  
...  
volume := 4/3 * pi * radius^3;
```

$4/3 * \pi$ can be evaluated at compile-time:

```
const pi: Double = 3.1415;  
var volume, radius: Double;  
...  
volume := 4.1888 * radius^3;
```

Constant Folding (2)

Not only applicable to *declared* constants:

```
x := 3;
```

```
y := x + 1;
```

```
x := x * 2;
```

can be optimized to

```
x := 3;
```

```
y := 4;
```

```
x := 6;
```

Constant Folding (3)

In general, *flow analysis* required:

```
x := 3;  
y := x + 1;  
while (x < z) begin  
    x := x * 2  
end
```

Unless z is known, we can only optimize to:

```
x := 3;  
y := 4;  
while (x < z) begin  
    x := x * 2  
end
```

Common Subexpression Elimination (1)

Idea: avoid evaluating the “same expression” more than once.

```
x1 := y1 + 7 * z + 42;
```

```
x2 := y2 + 7 * z + 42;
```

can be optimized to

```
t := 7 * z + 42;
```

```
x1 := y1 + t;
```

```
x2 := y2 + t;
```

Common subexpressions often appear in ***address computations*** in intermediate code.

Common Subexpression Elimination (2)

The expressions must not only be *syntactically* the same; they must also *mean* the same thing:

- Scope rules must be taken into account; consider Haskell-like `let`-expressions (i.e., functional code, *no* side effects):

```
let x = y * 17 in
  let y = 13 in
    let z = y * 17
```

The innermost `y * 17` *cannot* be replaced by `x`.

Common Subexpression Elimination (3)

- Side effects must be taken into account (flow analysis):

```
x := y * 17 + 3;
```

```
y := y + 1;
```

```
z := y * 17 + 3;
```

Here, the two instances of $y * 17 + 3$ do **not** compute the same value.

Indeed, the expressions themselves could have side effects (C-like increment operator):

```
x := y++ * 17 + 3;
```

```
z := y++ * 17 + 3;
```


Algebraic Identities (1)

Algebraic identities can be exploited to:

- simplify expressions: $1 * x - 0 \Rightarrow x$
- expose further opportunities for e.g. common subexpression evaluation:

```
x := (2 + z) * i;
```

```
y := (z + 2) * j;
```

can be transformed into

```
t := z + 2;
```

```
x := t * i;
```

```
y := t * j;
```

Algebraic Identities (2)

However, standard algebraic identities do not always hold!

Is it safe to assume that $x + (y + z)$ has the same meaning as $(x + y) + z$?

- Not if overflow/underflow is **trapped**: if x and y are large positive numbers, and z is a large negative number, then $(x + y) + z$ might result in a trap, while $x + (y + z)$ doesn't.
- **Floating point addition** is not associative!

Copy Propagation (1)

Idea: After an assignment that ***copies*** a value, like $x := y$ (often result of earlier optimization), use y in place of x wherever possible:

```
x := y;
```

```
v := x * 17;
```

```
w := x + 19;
```

can be transformed to

```
x := y;
```

```
v := y * 17;
```

```
w := y + 19;
```

Copy Propagation (2)

It may then turn out that the assigned variable is *never used again*. In that case, the assignment is *dead code* and can be eliminated.

```
x := y;  
v := y * 17;  
w := y + 19;
```

can be optimized to

```
v := y * 17;  
w := y + 19;
```

if x is never used again.

Dead Code Elimination (1)

Idea: It may be possible to **statically** determine that certain parts of the code

- will never be reached
- will not have any effect

The former is called **unreachable** code, the latter **dead** code.

Sometimes unreachable code is also referred to as dead code.

Either way, both are examples of **useless** code that can be **removed** without changing the meaning of the program.

Dead Code Elimination (2)

Consider the following Java fragment:

```
debug = false;
...
if (debug) {
    System.out.println("Got here!");
}
```

After constant folding, we have

```
if (false) {
    System.out.println("Got here!");
}
```

and the print statement is manifestly unreachable.

Dead Code Elimination (3)

In the copy propagation example, we saw that an assignment like

$$x := y;$$

could be removed if x is never used again as it has no effect and thus is dead code.

However, care needed: even if the assigned variable is never used, execution of the assignment statement itself might have an effect, meaning it **cannot** be removed (in its entirety):

$$x := y++;$$

Strength Reduction (1)

Idea: replace “expensive” operations by cheaper ones. Simple examples:

- Addition and shifting might be cheaper than multiplication:

$$5 * x \Rightarrow x \ll 2 + x$$

- Multiplication might be cheaper than exponentiation:

$$x^2 \Rightarrow x * x$$

$$z := x^5$$

$$\Rightarrow x2 := x * x; z := x2 * x2 * x$$

Only applies when **known** integral power.

Strength Reduction (2)

A loop may have a number of *induction variables* that remain in *lock step*:

```
i := 10;
while (i > 0) do begin
    i := i - 1;
    t := 4 * i;
    a[i] := b[t]
end
```

Here, *i* and *t* are induction variables.

Strength Reduction (3)

All that is going on is that t decreases by 4 each time round the loop. We can rephrase as follows:

```
i := 10;  
t := 4 * i;  
while (i > 0) do begin  
    i := i - 1;  
    t := t - 4;  
    a[i] := b[t]  
end
```

An potentially expensive multiplication has been replaced by a subtraction *inside* a loop.

Code Motion (1)

Idea: code that is *loop invariant*, i.e. evaluate to the same value at each loop iteration, should be moved outside the loop.

```
for (i := 0; i <= m - 1; i++) do
    for (j := 0; j <= n - 1; j++) do
        x := x + a[i * 10 + j]
```

- $m - 1$ and $n - 1$ invariant in the outer loop
- $i * 10$ invariant in the inner loop.

Code Motion (2)

Thus we can transform to:

```
t1 := m - 1;
t2 := n - 1;
for (i := 0; i <= t1; i++) do begin
    t3 := i * 10;
    for (j := 0; j <= t2; j++) do
        x := x + a[t3 + j]
    end
end
```

Array address computations and bounds checks often introduce loop invariant code fragments.

Code Motion (3)

Of course, we have to be careful if there are side effects. Consider:

```
for (i := 0; i < n; i++) do  
    x := x + f(17);
```

The function call $f(17)$ might look like loop invariant code at a first glance, but it could have side effects, in which case it is wrong to move it out of the loop:

```
f(n) = begin z := z + n; return z end;
```

•
•
•

Loop Unrolling (1)

As loops carry certain overheads (evaluation of loop condition, jumps), it can be beneficial to unroll loops that are known to be short. Consider:

```
for (i := 0; i < 5; i++) do  
    a[i] := b[4 - i] * 2^i;
```

Loop unrolling yields:

```
a[0] := b[4 - 0] * 2^0;  
a[1] := b[4 - 1] * 2^1;  
a[2] := b[4 - 2] * 2^2;  
a[3] := b[4 - 3] * 2^3;  
a[4] := b[4 - 4] * 2^4;
```

Loop Unrolling (2)

The resulting code can often be further improved; e.g. by constant folding:

```
a[0] := b[4] * 1;  
a[1] := b[3] * 2;  
a[2] := b[2] * 4;  
a[3] := b[1] * 8;  
a[4] := b[0] * 16;
```

Loop Unrolling (3)

Caveats:

- Loop unrolling can cause the code to grow considerably: space vs. time trade off.
- Impact of cache memories:
 - The instructions for a short loop will fit into the instruction cache and can thus be fetched again very quickly for each iteration.
 - Each instruction for an unrolled loop has to be fetched from main memory.

Loop Unrolling (4)

Loops where the bounds are statically unknown can sometimes still be partially unrolled:

```
for (i := 0; i < n; i++) do  
    a[i] := b[i] + c[i];
```

can for example be transformed into (integer div.):

```
for (i := 0; i < (n/2)*2; i := i+2) do begin  
    a[i] := b[i] + c[i];  
    a[i + 1] := b[i + 1] + c[i + 1]  
end;  
if (i < n) then begin  
    a[i] := b[i] + c[i];  
    i++  
end;
```

Loop Unrolling (5)

Benefits:

- Number of iterations reduced (here, roughly halved).
- Increased size of loop body may open up for further improvements; e.g. constant folding, CSE, strength reduction as discussed earlier (in particular for index address calculations).

Inlining (1)

Idea: Avoid overhead of function/procedure call by instantiating the body with the actual parameters and copying the result to the call site.

Also called *procedure integration*.

- Inlined procedures/functions should be *small*, or size of code might blow up!
- Careful with *recursion*! Otherwise the *compiler* might get stuck in a loop.
- Can make sense to unfold recursive procedures/functions a few times: similar to loop unrolling.

Inlining (2)

```
fun f (x: Integer): Integer =  
  begin  
    return (x + 17) * 123  
  end
```

...

```
x := f(a + 3);  
y := f(x * 3);
```

Inlining would result in the last fragment becoming:

```
x := ((a + 3) + 17) * 123;  
y := ((x * 3) + 17) * 123;
```

Inlining (3)

Consider:

```
fun fib (x : Integer) : Integer =  
begin  
    return (x < 2 ? x : fib(x-1) + fib(x-2))  
end
```

Recursion! Care needed!

If we blindly inline `fib` everywhere just because it initially looks small, compiler will get stuck in a loop (exhausting the memory eventually)!

•
•
•

Interaction among Optimizations (1)

One optimization might generate opportunities for other optimizations:

```
const level: Integer = 4;
const debugging: Boolean = true;
func debug(severity: Integer) =
begin
    return debugging && severity > level
end
...
x := 10;
if debug(3) then begin
    print "Oops! Well, got here.";
    x := x + 1
end;
y := x + 10;
```

•
•
•

Interaction among Optimizations (2)

Inlining yields:

```
const level: Integer = 4;
const debugging: Boolean = true;
...
x := 10;
if debugging && 3 > 4 then begin
    print "Oops! Well, got here.";
    x := x + 1
end;
y := x + 10;
```

Interaction among Optimizations (3)

Constant folding yields:

```
x := 10;  
if false then begin  
    print "Oops! Well, got here.";  
    x := x + 1  
end;  
y := x + 10;
```


•
•
•

Interaction among Optimizations (4)

Dead (unreachable) code elimination yields:

```
x := 10;
```

```
y := x + 10;
```

And now we can do further constant folding!

```
x := 10;
```

```
y := 20;
```

And then, if `x` never used again, more dead code elimination!

```
y := 20;
```

Interaction among Optimizations (5)

- In general hard to pick a “best” order among the optimizations.
- Compilers often carry out optimizations iteratively until no further improvements can be made.