Some Basic Mathematics for Machine Learning

Angela J. Yu <ajyu@ucsd.edu>

1 Probability Theory

See sections 2.1–2.3 of David MacKay's book: www.inference.phy.cam.ac.uk/mackay/itila/book.html

The probability a discrete variable A takes value a is: $0 \leq P(A{=}a) \leq 1$ — Probability

Probabilities of alternative outcomes add: $P(A{\in}\{a, a'\}) = P(A{=}a) + P(A{=}a')$ — Alternatives

The probabilities of all outcomes must sum to one: $\sum_{\text{all possible } a} P(A{=}a) = 1$ — Normalization

$P(A{=}a, B{=}b)$ is the joint probability that both $A{=}a$ and $B{=}b$ occur. — Joint Probability

Variables can be "summed out" of joint distributions: — Marginalization

$$P(A{=}a) = \sum_{\text{all possible } b} P(A{=}a, B{=}b)$$

$P(A{=}a|B{=}b)$ is the probability $A{=}a$ occurs given the knowledge $B{=}b$. — Conditional Probability

$P(A{=}a, B{=}b) = P(A{=}a)\, P(B{=}b|A{=}a) = P(B{=}b)\, P(A{=}a|B{=}b)$ — Product Rule

Bayes rule can be derived from the above: — Bayes Rule

$$P(A{=}a|B{=}b, \mathcal{H}) = \frac{P(B{=}b|A{=}a, \mathcal{H})\, P(A{=}a|\mathcal{H})}{P(B{=}b|\mathcal{H})} \propto P(A{=}a, B{=}b|\mathcal{H})$$

Marginalizing over all possible a gives the **evidence** or **normalizing constant**: — Normalizing Constant

$$\sum_a P(A{=}a, B{=}b|\mathcal{H}) = P(B{=}b|\mathcal{H})$$

The following hold, for all a and b, **if and only if A and B are independent**: — Independence

$$\begin{aligned} P(A{=}a|B{=}b) &= P(A{=}a) \\ P(B{=}b|A{=}a) &= P(B{=}b) \\ P(A{=}a, B{=}b) &= P(A{=}a)\, P(B{=}b)\,. \end{aligned}$$

All the above theory basically still applies to continuous variables if the sums are converted into integrals[1]. The probability that X lies between x and $x{+}\mathrm{d}x$ is $p(x)\,\mathrm{d}x$, where $p(x)$ is a *probability density function* with range $[0, \infty]$. — Continuous Variables

$$P(x_1{<}X{<}x_2) = \int_{x_1}^{x_2} p(x)\,\mathrm{d}x\,, \quad \int_{-\infty}^{\infty} p(x)\,\mathrm{d}x = 1 \;\text{ and }\; p(x) = \int_{-\infty}^{\infty} p(x, y)\,\mathrm{d}y.$$

The expectation or mean under a probability distribution is: — Expectation

$$\langle f(a)\rangle = \sum_a P(A{=}a)\, f(a) \quad \text{or} \quad \langle f(x)\rangle = \int_{-\infty}^{\infty} p(x)\, f(x)\mathrm{d}x$$

[1] Integrals are the equivalent of sums for continuous variables, e.g $\sum_{i=1}^{n} f(x_i)\Delta x$ becomes the integral $\int_a^b f(x)\mathrm{d}x$ in the limit $\Delta x \to 0$, $n \to \infty$, where $\Delta x = \frac{b-a}{n}$ and $x_i = a + i\Delta x$.

2 Linear Algebra

This complements Sam Roweis's "Matrix Identities": `www.cs.toronto.edu/~roweis/notes/matrixid.pdf`

Scalars are individual numbers, vectors are columns of numbers, matrices are rectangular grids of numbers, eg:

$$x = 3.4, \qquad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}, \qquad A = \begin{pmatrix} A_{11} & A_{12} & \cdots & A_{1n} \\ A_{21} & A_{22} & \cdots & A_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ A_{m1} & A_{m2} & \cdots & A_{mn} \end{pmatrix}$$

In the above example x is 1×1, $\mathbf{x}$ is $n \times 1$ and A is $m \times n$. — Dimensions

The transpose operator, $^\top$ ($'$ in Matlab), swaps the rows and columns: — Transpose

$$x^\top = x, \qquad \mathbf{x}^\top = \begin{pmatrix} x_1 & x_2 & \cdots & x_n \end{pmatrix}, \qquad \left(A^\top\right)_{ij} = A_{ji}$$

Quantities whose inner dimensions match may be "multiplied" by summing over this index. The outer dimensions give the dimensions of the answer. — Multiplication

$$A\mathbf{x} \text{ has elements } (A\mathbf{x})_i = \sum_{j=1}^{n} A_{ij}\mathbf{x}_j \quad \text{and} \quad (AA^\top)_{ij} = \sum_{k=1}^{n} A_{ik}\left(A^\top\right)_{kj} = \sum_{k=1}^{n} A_{ik}A_{jk}$$

All the following are allowed (the dimensions of the answer are also shown): — Check Dimensions

$\mathbf{x}^\top\mathbf{x}$	$\mathbf{x}\mathbf{x}^\top$	$A\mathbf{x}$	$AA^\top$	$A^\top A$	$\mathbf{x}^\top A\mathbf{x}$
1×1	$n \times n$	$m \times 1$	$m \times m$	$n \times n$	1×1
scalar	matrix	vector	matrix	matrix	scalar

,

while $\mathbf{x}\mathbf{x}$, AA and $\mathbf{x}A$ *do not make sense* for $m \neq n \neq 1$. Can you see why?

An exception to the above rule is that we may write: xA. Every element of the matrix A is multiplied by the scalar x. — Multiplication by Scalar

Simple and valid manipulations: — Basic Properties

$$(AB)C = A(BC) \qquad A(B+C) = AB + AC \qquad (A+B)^\top = A^\top + B^\top \qquad (AB)^\top = B^\top A^\top$$

Note that $AB \neq BA$ in general.

2.1 Square Matrices

A square matrix has equal number of rows and columns, e.g. B with dimensions $n \times n$.

A diagonal matrix is a square matrix with all off-diagonal elements being zero: $B_{ij} = 0$ if $i \neq j$. — Diagonal Matrix

The identity matrix is a diagonal matrix with all diagonal elements equal to one. — Identity Matrix

$$\text{"}\mathbb{I} \text{ is the identity matrix"} \quad \Leftrightarrow \quad \mathbb{I}_{ij} = 0 \text{ if } i \neq j \text{ and } \mathbb{I}_{ii} = 1 \ \ \forall i$$

The identity matrix leaves vectors and matrices unchanged upoon multiplication.

$$\mathbb{I}\mathbf{x} = \mathbf{x} \qquad \mathbb{I}B = B = B\mathbb{I} \qquad \mathbf{x}^\top\mathbb{I} = \mathbf{x}^\top$$

Some square matrices have inverses: Inverse

$$B^{-1}B = BB^{-1} = \mathbb{I} \qquad \left(B^{-1}\right)^{-1} = B\,,$$

which have the additional properties:

$$(BC)^{-1} = C^{-1}B^{-1} \qquad \left(B^{-1}\right)^{\top} = \left(B^{\top}\right)^{-1}$$

Linear simultaneous equations could be solved this way: Solving Linear Equations

$$\text{if } B\mathbf{x} = \mathbf{y} \text{ then } \mathbf{x} = B^{-1}\mathbf{y}$$

Some other commonly used matrix definitions include:

$$B_{ij} = B_{ji} \Leftrightarrow \text{“}B \text{ is symmetric”} \qquad \text{Symmetry}$$

$$\text{Trace}(B) = \text{Tr}(B) = \sum_{i=1}^{n} B_{ii} = \text{“sum of diagonal elements”} \qquad \text{Trace}$$

Cyclic permutations are allowed inside trace. Trace of a scalar is a scalar: A Trace Trick

$$\text{Tr}(BCD) = \text{Tr}(DBC) = \text{Tr}(CDB) \qquad \mathbf{x}^{\top}B\mathbf{x} = \text{Tr}(\mathbf{x}^{\top}B\mathbf{x}) = \text{Tr}(\mathbf{x}\mathbf{x}^{\top}B)$$

The determinant[2] is written $\text{Det}(B)$ or $|B|$. It is a scalar regardless of n. Determinant

$$|BC| = |B||C|\,, \qquad |x| = x\,, \qquad |xB| = x^n|B|\,, \qquad \left|B^{-1}\right| = \frac{1}{|B|}\,.$$

It *determines* if B can be inverted: $|B| = 0 \Rightarrow B^{-1}$ undefined. If the vector to every point of a shape is pre-multiplied by B then the shape's area or volume increases by a factor of $|B|$. It also appears in the normalising constant of a Gaussian. For a diagonal matrix the volume scaling factor is simply the product of the diagonal elements. In general the determinant is the product of the eigenvalues.

$$B\mathbf{v}^{(i)} = \lambda^{(i)}\mathbf{v}^{(i)} \Leftrightarrow \text{“}\lambda^{(i)} \text{ is an eigenvalue of } B \text{ with eigenvector } \mathbf{v}^{(i)}\text{”} \qquad \text{Eigenvalue, Eigenvector}$$

$$|B| = \prod_i \lambda^{(i)} \qquad \text{Trace}(B) = \sum_i \lambda^{(i)}$$

If B is real and symmetric (eg a covariance matrix), the eigenvectors are orthogonal (perpendicular) and so form a basis (can be used as axes).

[2]This section is only intended to give you a flavour so you understand other references and Sam's crib sheet. More detailed history and overview is here: `http://www.wikipedia.org/wiki/Determinant`

3 Differentiation

Gradient

The gradient of a straight line $y=mx+c$ is a constant: $y' = \frac{y(x+\Delta x)-y(x)}{\Delta x} = m$.

Differentiation

Many functions look like straight lines over a small enough range. The gradient of this line, the derivative, is not constant, but a new function:

$$y'(x) = \frac{\mathrm{d}y}{\mathrm{d}x} = \lim_{\Delta x \to 0} \frac{y(x+\Delta x) - y(x)}{\Delta x}\,, \qquad \text{which could be differentiated again:} \quad y'' = \frac{\mathrm{d}^2 y}{\mathrm{d}x^2} = \frac{\mathrm{d}y'}{\mathrm{d}x}$$

Standard Derivatives

The following results are well known (c is a constant):

$f(x):$	c	cx	cx^n	$\log_e(x)$	e^x
$f'(x):$	0	c	cnx^{n-1}	$1/x$	e^x

Optimization

At a maximum or minimum the function is rising on one side and falling on the other. In between the gradient must reach zero somewhere. Therefore

$$\text{maxima and minima satisfy: } \frac{\mathrm{d}f(x)}{\mathrm{d}x} = 0 \qquad \text{or} \qquad \frac{\mathrm{d}f(\mathbf{x})}{\mathrm{d}\mathbf{x}} = \mathbf{0} \;\Leftrightarrow\; \frac{\mathrm{d}f(\mathbf{x})}{\mathrm{d}x_i} = 0 \quad \forall i$$

If we can't solve this analytically, we can evolve our variable x, or variables $\mathbf{x}$, on a computer using gradient information until we find a place where the gradient is zero. But there's no guarantee that we would find all maxima and minima.

Approximation

A function may be approximated by a straight line[3] about any point a.

$$f(a+x) \approx f(a) + xf'(a)\,, \qquad \text{e.g.: } \log(1+x) \approx \log(1+0) + x\frac{1}{1+0} = x$$

Linearity

The derivative operator is linear:

$$\frac{\mathrm{d}(f(x)+g(x))}{\mathrm{d}x} = \frac{\mathrm{d}f(x)}{\mathrm{d}x} + \frac{\mathrm{d}g(x)}{\mathrm{d}x}\,, \qquad \text{e.g.: } \frac{\mathrm{d}\left(x+\exp(x)\right)}{\mathrm{d}x} = 1+\exp(x).$$

Product Rule

Dealing with products is slightly more involved:

$$\frac{\mathrm{d}\left(u(x)v(x)\right)}{\mathrm{d}x} = v\frac{\mathrm{d}u}{\mathrm{d}x} + u\frac{\mathrm{d}v}{\mathrm{d}x}\,, \qquad \text{e.g.: } \frac{\mathrm{d}\left(x\cdot\exp(x)\right)}{\mathrm{d}x} = \exp(x) + x\exp(x).$$

Chain Rule

The "chain rule" $\frac{\mathrm{d}f(u)}{\mathrm{d}x} = \frac{\mathrm{d}u}{\mathrm{d}x}\frac{\mathrm{d}f(u)}{\mathrm{d}u}$, allows results to be combined.

$$\text{For example: } \frac{\mathrm{d}\exp\left(ay^m\right)}{\mathrm{d}y} = \frac{\mathrm{d}\left(ay^m\right)}{\mathrm{d}y} \cdot \frac{\mathrm{d}\exp\left(ay^m\right)}{\mathrm{d}\left(ay^m\right)} \quad \text{"with } u = ay^m\text{"}$$
$$= amy^{m-1} \cdot \exp\left(ay^m\right)$$

Exercise

Convince yourself of the following:

$$\frac{\mathrm{d}}{\mathrm{d}z}\left[\frac{1}{(b+cz)}\exp(az) + e\right] = \exp(az)\left(\frac{a}{b+cz} - \frac{c}{(b+cz)^2}\right)$$

Note that a, b, c and e are constants and $1/u = u^{-1}$. This might be hard if you haven't done differentiation (for a long time). You may want to grab a calculus textbook for a review.

[3] More accurate approximations can be made. Look up Taylor series.